

Deep Joint-Semantics Reconstructing Hashing for Large-Scale Unsupervised Cross-Modal Retrieval

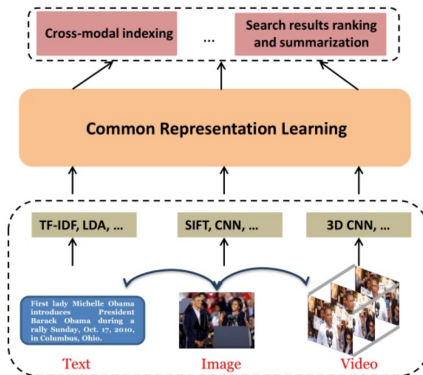
Oriol Vinyals¹ Zhisheng Zhong²

¹Key Laboratory of Machine Perception (MOE), School of EECS, Peking University

²Key Laboratory of Machine Perception (MOE), School of EECS, Peking University

Unsupervised Cross-Modal Retrieval

Cross-Modal Retrieval



Binary representation learning *Unsupervised Cross-Modal Retrieval*

Contributions

- Put forward affinity matrix.
- Reconstruct above jointsemantics, friendly for batch-wise training.
- Reach good result.

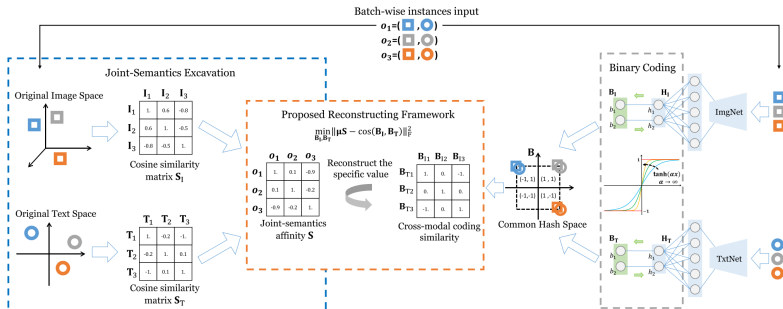


图: The pipeline of DJSRH

- m : Batch size;
- \mathcal{O} : $\{o_k = [I_k, T_k]\}_{k=1}^m$, include each image-text pair. Feature matrix of image and text are defined as $F_I \in \mathbb{R}^{m \times p_I}$ and $F_T \in \mathbb{R}^{m \times p_T}$;
- $B_I \in \{\pm 1\}^{m \times d}$ and $B_T \in \{\pm 1\}^{m \times d}$: Binary representation given out by ImgNet and TxtNet from the input I_k and T_k ;
- \hat{F}_I and \hat{F}_T : The normalized F_I and F_T , the cosine similarity matrices

$$S_I = \hat{F}_I \hat{F}_I^T \in [-1, +1]^{m \times m}$$

$$S_T = \hat{F}_T \hat{F}_T^T \in [-1, +1]^{m \times m}$$

Constructing Joint-Semantics Matrix

Laplacian constrains

$$\min_B \beta \text{Tr}(B^\top L_I B) + (1 - \beta) \text{Tr}(B^\top L_T B) \quad \text{s.t. } B \in \{\pm 1\}^{m \times d}$$

where

$$L_I = \text{diag}(S_I \mathbf{1}) - S_I$$

$$L_T = \text{diag}(S_T \mathbf{1}) - S_T$$

are Laplacian matrices.

Constructing Joint-Semantics Matrix

Joint-semantics Affinity Matrix

Define \mathcal{C} as combination function, then

$$S = \mathcal{C}(S_I, S_T) \in [-1, +1]^{m \times m}$$

Merge Img and Txt as

$$\tilde{S} = \beta S_I + (1 - \beta) S_T$$

Then

$$S = \mathcal{C}(S_I, S_T)$$

$$= (1 - \eta) \tilde{S} + \eta \frac{\tilde{S} \tilde{S}^\top}{m}$$

$$= (1 - \eta) [\beta S_I + (1 - \beta) S_T] + \frac{\eta}{m} [\beta^2 S_I S_I^\top + \beta(1 - \beta)(S_I S_T^\top + S_T S_I^\top)]$$

S_{ij} indicates the latent semantic similarity between o_i 和 o_j .

Reconstructing with Binary Codes

Object

$$\min_{B_I, B_T} \|\mu S - \cos(B_I, B_T)\|_F^2, \quad \text{s.t. } S = \mathcal{C}(S_I, S_T) \in [-1, +1]^{m \times m}$$

Laplacian constrains

$$\text{Tr}(B^T L B) = \sum_{i,j} S_{ij} \|B_i - B_j\|^2$$

Object with intra-modal influence

$$\begin{aligned} \min_{B_I, B_T} & \|\mu S - \cos(B_I, B_T)\|_F^2 + \lambda_1 \|\mu S - \cos(B_I, B_I)\|_F^2 \\ & + \lambda_2 \|\mu S - \cos(B_T, B_T)\|_F^2, \\ \text{s.t. } & S = \mathcal{C}(S_I, S_T) \in [-1, +1]^{m \times m}, \quad B_I, B_T \in \{-1, +1\}^{m \times d} \end{aligned}$$

Set $H \in \mathbb{R}^{m \times d}$ as the last layer of ImgNet and TxtNet without activate function, then

$$B = \text{sgn}(H) \in \{-1, +1\}^{m \times d}$$

Use the following instead,

$$B = \tanh(\alpha H) \in \{-1, +1\}^{m \times d}, \alpha \in \mathbb{R}^+$$

Algorithm 1 Deep Joint-Semantics Reconstructing Hashing

Input:

Training set $\{\mathbf{o}_k = [\mathbf{I}_k, \mathbf{T}_k]\}_{k=1}^n$ and their corresponding original features \mathbf{F}_I and \mathbf{F}_T ; ImgNet \mathcal{G}_{θ_I} and TxtNet \mathcal{G}_{θ_T} with θ_I and θ_T denoting the deep network parameters; batch size m ;

Output:

Hashing coding function $\varphi_I(x) = \text{sgn}(\mathcal{G}_{\theta_I}(x))$ for image input and $\varphi_T(x) = \text{sgn}(\mathcal{G}_{\theta_T}(x))$ for text input;

- 1: Initialize epoch $t = 0$;
 - 2: **repeat**
 - 3: $t = t + 1$; $\alpha = \sqrt{t}$;
 - 4: **for** $\lfloor \frac{n}{m} \rfloor$ iterations **do**
 - 5: Randomly sample a batch of instances from training set $\{\mathbf{o}_k = [\mathbf{I}_k, \mathbf{T}_k]\}_{k=1}^m$;
 - 6: Calculate the normalized $\hat{\mathbf{F}}_I, \hat{\mathbf{F}}_T$ and integrate the cosine matrices $\mathbf{S}_I = \hat{\mathbf{F}}_I \hat{\mathbf{F}}_I^\top, \mathbf{S}_T = \hat{\mathbf{F}}_T \hat{\mathbf{F}}_T^\top$ to the joint-semantics affinity \mathbf{S} with Equation (3);
 - 7: Forward propagate $\mathbf{H}_I = \mathcal{G}_{\theta_I}(\mathbf{I}), \mathbf{H}_T = \mathcal{G}_{\theta_T}(\mathbf{T})$;
 - 8: Hash coding with activation function (7) $\mathbf{B}_I = \tanh(\alpha \mathbf{H}_I), \mathbf{B}_T = \tanh(\alpha \mathbf{H}_T)$;
 - 9: Calculate the objective function (5), back propagate the gradients with the chain rule and update the whole parameters;
 - 10: **end for**
 - 11: **until** convergence
-

Algorithm

```
F_I = F.normalize(F_I)
S_I = F_I.mm(F_I.t())
S_I = S_I * 2 - 1

F_T = F.normalize(F_T)
S_T = F_T.mm(F_T.t())
S_T = S_T * 2 - 1

B_I = F.normalize(code_I)
B_T = F.normalize(code_T)

BI_BI = B_I.mm(B_I.t())
BT_BT = B_T.mm(B_T.t())
BI_BT = B_I.mm(B_T.t())

S_tilde = settings.BETA * S_I + (1 - settings.BETA) * S_T
S = (1 - settings.ETA) * S_tilde + settings.ETA * S_tilde.mm(S_tilde) / settings.BATCH_SIZE
S = S * settings.MU

loss1 = F.mse_loss(BI_BI, S)
loss2 = F.mse_loss(BI_BT, S)
loss3 = F.mse_loss(BT_BT, S)
loss = settings.LAMBDA1 * loss1 + 1 * loss2 + settings.LAMBDA2 * loss3
```

Experiment

| Task | Method | Wiki | | | | MIRFlickr | | | | NUS-WIDE | | | |
|-------------------|--------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|
| | | 16bits | 32bits | 64bits | 128bits | 16bits | 32bits | 64bits | 128bits | 16bits | 32bits | 64bits | 128bits |
| $I \rightarrow T$ | CVH | 0.179 | 0.162 | 0.153 | 0.149 | 0.606 | 0.599 | 0.596 | 0.598 | 0.372 | 0.362 | 0.406 | 0.390 |
| | IMH | 0.201 | 0.203 | 0.204 | 0.195 | 0.612 | 0.601 | 0.592 | 0.579 | 0.470 | 0.473 | 0.476 | 0.459 |
| | CMFH | 0.251 | 0.253 | 0.259 | 0.263 | 0.621 | 0.624 | 0.625 | 0.627 | 0.455 | 0.459 | 0.465 | 0.467 |
| | LSSH | 0.197 | 0.208 | 0.199 | 0.195 | 0.584 | 0.599 | 0.602 | 0.614 | 0.481 | 0.489 | 0.507 | 0.507 |
| | DBRC | 0.253 | 0.265 | 0.269 | 0.288 | 0.617 | 0.619 | 0.620 | 0.621 | 0.424 | 0.459 | 0.447 | 0.447 |
| | UDCMH | 0.309 | 0.318 | 0.329 | 0.346 | 0.689 | 0.698 | 0.714 | 0.717 | 0.511 | 0.519 | 0.524 | 0.558 |
| | DJSRH | 0.388 | 0.403 | 0.412 | 0.421 | 0.810 | 0.843 | 0.862 | 0.876 | 0.724 | 0.773 | 0.798 | 0.817 |
| $T \rightarrow I$ | CVH | 0.252 | 0.235 | 0.171 | 0.154 | 0.591 | 0.583 | 0.576 | 0.576 | 0.401 | 0.384 | 0.442 | 0.432 |
| | IMH | 0.467 | 0.478 | 0.453 | 0.456 | 0.603 | 0.595 | 0.589 | 0.580 | 0.478 | 0.483 | 0.472 | 0.462 |
| | CMFH | 0.595 | 0.601 | 0.616 | 0.622 | 0.642 | 0.662 | 0.676 | 0.685 | 0.529 | 0.577 | 0.614 | 0.645 |
| | LSSH | 0.569 | 0.593 | 0.593 | 0.595 | 0.637 | 0.659 | 0.659 | 0.672 | 0.577 | 0.617 | 0.642 | 0.663 |
| | DBRC | 0.574 | 0.588 | 0.598 | 0.599 | 0.618 | 0.626 | 0.626 | 0.628 | 0.455 | 0.459 | 0.468 | 0.473 |
| | UDCMH | 0.622 | 0.633 | 0.645 | 0.658 | 0.692 | 0.704 | 0.718 | 0.733 | 0.637 | 0.653 | 0.695 | 0.716 |
| | DJSRH | 0.611 | 0.635 | 0.646 | 0.658 | 0.786 | 0.822 | 0.835 | 0.847 | 0.712 | 0.744 | 0.771 | 0.789 |

图: mAP@50

MIRFlickr, CODELEN = 64

- mAP of Image to Text: 0.865
- mAP of Text to Image: 0.853